

Dissecting Positional Encoding in Encoder-Decoder MT:

A Cross-Lingual Study via Adaptive RoPE

Nilesh Mishra

Department of Computer Science and Engineering

Your Institution

thenileshmishra@gmail.com

Abstract—Rotary Position Embedding (RoPE) has become the dominant positional encoding in neural machine translation, yet its core assumption — that all frequency components contribute equally — remains empirically untested in encoder-decoder architectures. We introduce AdaptiveRoPE, a minimal extension that augments RoPE with learnable per-head frequency gates and phase offsets. Initialised to recover standard RoPE, these parameters reveal which frequencies the model actually learns to use. We evaluate on three language pairs (En–De with 3 seeds, Hi–En and Bn–En) against 5 baselines (RoPE, Sinusoidal, ALiBi, and two ablations). Our findings are threefold: (1) trained gates consistently converge below one (mean ≈ 0.90), confirming a low-frequency preference; (2) the optimal adaptation is language-dependent: phase offsets are critical for Hindi–En (+21% BLEU over RoPE) while frequency gates matter more for Bengali–En; (3) RoPE achieves the highest in-domain BLEU, but ALiBi degrades far less on out-of-distribution lengths (−47% vs. −97% at 201–400 words), revealing a fundamental accuracy–extrapolation trade-off. These results suggest that “one-size-fits-all” positional encodings are suboptimal and motivate language-adaptive designs.

Index Terms—rotary position embedding, machine translation, positional encoding, frequency analysis, cross-lingual study, encoder-decoder transformer

I. INTRODUCTION

Position-aware attention is the backbone of sequence-to-sequence models. Since the original sinusoidal encodings [1], a succession of relative and rotary approaches have improved how Transformers encode position [2]–[4]. Rotary Position Embedding (RoPE) [5] has emerged as the de facto standard: it rotates query and key vectors in 2-D frequency planes, making the attention score depend only on relative position while preserving the key–value cache and compilation compatibility. RoPE is now used in LLaMA [6], Mistral [7], and state-of-the-art MT systems.

Despite its ubiquity, a critical assumption of RoPE has never been validated in encoder-decoder MT: *all frequency components receive equal weight*. RoPE assigns each dimension pair $(2i, 2i+1)$ a fixed rotation frequency $\omega_i = 10000^{-2i/d}$, with no mechanism to learn that some frequencies are more useful than others. Recent work on decoder-only LLMs shows that high-frequency RoPE dimensions can be pruned with minimal perplexity loss [8], suggesting the uniform weighting

is wasteful. However, encoder-decoder MT differs fundamentally: positional information must serve a bidirectional encoder, a causal decoder, *and* cross-attention between them. Whether the same frequency preference holds — and whether it varies across languages — is unknown.

Research questions. This paper asks:

- RQ1** Do encoder-decoder MT models learn to suppress or amplify specific RoPE frequency components?
- RQ2** Do the learned adaptations generalise across language pairs, or are positional encoding requirements language-specific?
- RQ3** How do RoPE variants compare to alternatives (Sinusoidal, ALiBi) in terms of in-domain accuracy, length generalisation, and efficiency?

Approach. We design AdaptiveRoPE, a differentiable probe that adds learnable per-head gates $\mathbf{g} \in \mathbb{R}^{H \times F}$ and phase offsets $\phi \in \mathbb{R}^{H \times F}$ to RoPE’s rotation angles (Section III). Initialised at standard RoPE ($\mathbf{g} = \mathbf{1}$, $\phi = \mathbf{0}$), any divergence reflects a genuine learned preference. Two ablations — **GatesOnly** (phases fixed at zero) and **PhasesOnly** (gates fixed at one) — disentangle the contribution of each component.

We train on WMT14 En–De (3 seeds), Samanantar Hi–En, and Samanantar Bn–En with a 47M-parameter encoder-decoder Transformer. We evaluate 6 positional encoding methods with BLEU, chrF, TER, and attention entropy; test statistical significance; and measure length generalisation on 8 buckets up to 400 source words.

Key findings.

- Gates converge to ≈ 0.90 (Section VI-A), confirming systematic low-frequency preference in encoder-decoder MT.
- *Cross-lingual divergence*: PhasesOnly outperforms full AdaptiveRoPE on Hi–En (+21% BLEU over RoPE), while AdaptiveRoPE itself wins on Bn–En (+0.8% over RoPE). This reveals that the optimal PE strategy is language-dependent (Section V-B).
- RoPE achieves the highest in-domain BLEU, but ALiBi generalises better to OOD lengths (−47% vs. −97% degradation, Section V-D).

- Attention entropy analysis shows ALiBi produces more uniform attention distributions (higher entropy), explaining its robustness to long inputs (Section V-E).

II. RELATED WORK

Positional encodings. Vaswani et al. [1] introduced sinusoidal absolute encodings. Shaw et al. [2] proposed relative position representations. ALiBi [4] replaces explicit position encodings with a linear distance-based bias in attention logits, achieving strong length extrapolation. RoPE [5] uses rotation matrices to make dot products depend on relative position; it dominates modern architectures due to cache efficiency and compilation-friendly kernels.

Analysing RoPE. Men et al. [8] study which RoPE frequencies are used in decoder-only Gemma-7B and find that the highest frequencies can be truncated (p-RoPE). Our work differs in four ways: (i) we study *encoder-decoder* MT, not decoder-only LLMs; (ii) we use a *differentiable probe* rather than post-hoc ablation; (iii) we disentangle *gates* from *phases* via ablations; (iv) we evaluate across *multiple language pairs* (En-De, Hi-En, Bn-En).

Learnable RoPE variants. YaRN [9] and LongRoPE [10] learn frequency scalings for context-window extension. These target length generalisation; we target understanding what frequencies are needed at training-time lengths, with ablations that isolate mechanisms.

Cross-lingual PE studies. Most PE research evaluates on English or En-De [4], [5], [8]. To our knowledge, no prior work systematically compares PE requirements across morphologically diverse language pairs with controlled architecture and training.

III. METHOD

A. Rotary Position Embedding (RoPE)

For a query vector $\mathbf{q} \in \mathbb{R}^d$ at position p , RoPE applies a block-diagonal rotation to each pair $(2i, 2i+1)$:

$$\begin{bmatrix} q'_{2i} \\ q'_{2i+1} \end{bmatrix} = \begin{bmatrix} \cos \theta_p^{(i)} & -\sin \theta_p^{(i)} \\ \sin \theta_p^{(i)} & \cos \theta_p^{(i)} \end{bmatrix} \begin{bmatrix} q_{2i} \\ q_{2i+1} \end{bmatrix}, \quad (1)$$

where $\theta_p^{(i)} = p \cdot \omega_i$ and $\omega_i = 10000^{-2i/d}$. The same rotation is applied to keys, so the inner product depends only on relative offset $p - p'$ [5].

B. Adaptive RoPE

We augment the rotation angle with head-specific learnable gates $\mathbf{g}$ and phase offsets ϕ :

$$\theta_p^{(h,i)} = p \cdot \omega_i \cdot g^{(h,i)} + \phi^{(h,i)}, \quad (2)$$

where $h \in \{1, \dots, H\}$ indexes the attention head and $i \in \{0, \dots, F-1\}$ indexes the frequency ($F = d_{\text{head}}/2$). Separate parameters are maintained for queries and keys: $\mathbf{g}_q, \mathbf{g}_k, \phi_q, \phi_k \in \mathbb{R}^{H \times F}$.

Initialisation. $\mathbf{g} = \mathbf{1}$, $\phi = \mathbf{0}$. This makes AdaptiveRoPE *identical* to standard RoPE at initialisation.

Interpretation. $g^{(h,i)} < 1$ suppresses frequency i ; $g^{(h,i)} > 1$ amplifies it. $\phi^{(h,i)}$ shifts the rotation angle independently of frequency.

Overhead. For $H = 6$, $F = 32$, AdaptiveRoPE adds $4 \times 6 \times 32 = 768$ scalars per layer ($< 0.002\%$ of 47M parameters).

C. Ablations

To isolate mechanisms we define two ablations: **GatesOnly** has learnable $\mathbf{g}$ with $\phi = \mathbf{0}$; **PhasesOnly** has learnable ϕ with $\mathbf{g} = \mathbf{1}$. Comparing these to full AdaptiveRoPE and standard RoPE reveals whether frequency scaling or phase shifting drives performance on each language pair.

IV. EXPERIMENTAL SETUP

A. Datasets

WMT14 English-German [11]: 4.5M training pairs, newstest2014 test (3,003 sentences). Tokenisation: Helsinki-NLP/opus-mt-en-de SentencePiece (58K vocab).

Samanantar Hi-En and Bn-En [12]: 1M and 450K training pairs respectively. Test: 3,000 sentences sampled from official test splits. Tokenisation: ai4bharat/IndicBART (64K vocab) [13].

B. Model and Training

Standard encoder-decoder Transformer: 6+6 layers, $d_{\text{model}} = 384$, $H = 6$ heads, $d_{\text{ff}} = 1536$, dropout 0.1, $\approx 47\text{M}$ parameters. Embeddings tied with output projection.

Trained for 25K–30K steps, effective batch size 512, AdamW [14] with $\text{lr } 10^{-3}$, cosine schedule, 5% warmup, gradient clipping 1.0, bf16 mixed precision, torch.compile, label smoothing $\varepsilon = 0.1$.

En-De uses seeds 42, 43, 44; Hi-En and Bn-En use seed 42.

C. Evaluation Metrics

BLEU [15], **chrF** [16], and **TER** (Translation Edit Rate) via SacreBLEU [17]. Decoding: greedy (main results) and beam search ($b = 5$) for comparison.

Attention analysis: mean entropy and sink ratio (fraction of attention mass on the first token) per layer, computed on 320 validation samples.

Length generalisation: 8 source-length buckets (1–20, 21–40, ..., 201–400 words). PE buffers extended to 420 tokens where needed.

Statistical testing: paired t -tests for multi-seed comparisons (En-De); bootstrap resampling with 1,000 iterations for single-seed pairs.

V. RESULTS

A. Translation Quality

Table II reports BLEU, chrF, TER, and inference speed for all methods across three language pairs.

En-De. RoPE achieves the highest mean BLEU (20.83 ± 0.03), followed by PhasesOnly (20.75 ± 0.05), AdaptiveRoPE

TABLE I
PAIRWISE STATISTICAL SIGNIFICANCE ON EN-DE BLEU (WELCH'S
 t -TEST). ** $p < 0.01$, * $p < 0.05$.

	RoPE	ARoPE	Gates	Phases	Sinus	ALiBi
RoPE	—	.001**	.002**	.034*	<.001**	N/A
ARoPE	—	—	.835	.162	<.001**	N/A
Gates	—	—	—	.142	<.001**	N/A
Phases	—	—	—	—	<.001**	N/A
Sinus	—	—	—	—	—	N/A

(20.69 ± 0.05), and GatesOnly (20.68 ± 0.08). All RoPE-based methods are tightly clustered within 0.15 BLEU, but the differences are statistically significant due to low inter-seed variance: RoPE vs. AdaptiveRoPE ($p = 0.001$), RoPE vs. GatesOnly ($p = 0.002$), and RoPE vs. PhasesOnly ($p = 0.034$). Sinusoidal (20.18 ± 0.15) and ALiBi (19.60) lag by 0.6–1.2 BLEU, both highly significant ($p < 0.001$).

Hi-En. The ranking reverses dramatically: PhasesOnly achieves 9.42 BLEU, outperforming AdaptiveRoPE (9.08), GatesOnly (8.37), and RoPE (7.78) — a **21% improvement** of PhasesOnly over RoPE. This is our strongest cross-lingual finding: phase offsets are critical for Hindi–English translation. Notably, full AdaptiveRoPE (9.08) underperforms PhasesOnly (9.42), suggesting that the gates component interferes with the beneficial phase adaptation on this language pair.

Bn-En. AdaptiveRoPE wins (18.01), followed by RoPE (17.87) and GatesOnly (17.79). PhasesOnly drops to 17.06, the lowest among RoPE variants and 0.8 BLEU below RoPE. This is the *opposite* of Hi-En: gates (frequency scaling) matter more than phases for Bengali. The full AdaptiveRoPE outperforms both ablations, suggesting gates and phases are complementary for this language pair.

Speed. ALiBi is fastest (55.6s) because it avoids RoPE's rotary computation. Sinusoidal is similarly fast (54.7s). RoPE-based methods incur a small overhead (≈ 12 – 29%) due to rotary kernel calls. AdaptiveRoPE's overhead is negligible (≈ 2 s) relative to standard RoPE.

Statistical significance. Table I reports pairwise p -values from Welch's t -tests on En-De BLEU scores. RoPE is significantly better than all other methods except PhasesOnly (marginal, $p = 0.034$). Among RoPE variants, AdaptiveRoPE and GatesOnly are not distinguishable ($p = 0.835$), nor are AdaptiveRoPE and PhasesOnly ($p = 0.162$). This confirms that while RoPE maintains a small but consistent edge, the ablations cluster together.

B. Cross-Lingual Ablations

The ablation results reveal language-dependent PE requirements (Figure 1).

Morphologically rich + free word order (Hindi). PhasesOnly > AdaptiveRoPE > GatesOnly > RoPE. Phase offsets — which shift rotations without changing frequency scale — are critical for capturing Hindi's flexible syntax and rich morphology. RoPE's fixed phases are a constraint the model strongly prefers to relax.




Fig. 1. BLEU comparison across language pairs. Error bars show std over seeds (En-De only).

Synthetic + relatively fixed order (Bengali). AdaptiveRoPE > RoPE $\approx$ GatesOnly > PhasesOnly. Frequency scaling (gates) is more important than phase shifts. Bengali's more regular word order and agglutinative morphology benefit from adaptive frequency weighting.

Synthetic + rigid order (German). All RoPE variants cluster within 0.15 BLEU (20.68–20.83). Neither component dominates, suggesting RoPE's defaults are near-optimal for Germanic syntax.

These patterns suggest that *positional encoding requirements are not universal*: morphological richness and word-order flexibility determine whether phase or frequency adaptation is more valuable.

C. Per-Length Analysis

Table III breaks BLEU down by source sentence length on En-De. All methods perform best on mid-length sentences (11–50 words), which dominate the training distribution. On very short sentences (1–10 words), RoPE leads by 0.3–0.7 BLEU over ablations, suggesting that its uniform frequency spectrum is well-suited for brief inputs. On long sentences (51+ words, 21 samples), RoPE and AdaptiveRoPE tie (≈ 22.6), with ALiBi close behind (22.2).

D. Length Generalisation

Table IV reports BLEU degradation on OOD source lengths (En-De). ALiBi shows dramatically better extrapolation: only -47% degradation at 201–400 words, compared to -97% for RoPE and -98% for AdaptiveRoPE variants. Sinusoidal also generalises better (-96%) than RoPE-based methods.

Among RoPE variants, differences are small: GatesOnly (-31.8% at 81–120 words) marginally outperforms RoPE (-36.7%), but all converge to near-zero BLEU at 201–400 words. The learned adaptations provide insufficient extrapolation benefit to close the gap with ALiBi.

This reveals a fundamental **accuracy–extrapolation trade-off**: RoPE achieves the highest in-domain BLEU but degrades most rapidly on long sequences. ALiBi sacrifices 1.2 BLEU in-domain for robust extrapolation. For deployment with variable input lengths, the choice between RoPE and ALiBi depends on the expected length distribution.

E. Attention Pattern Analysis

Table V reports mean attention entropy and sink ratio. ALiBi exhibits substantially higher entropy (2.233) than all

TABLE II

TRANSLATION QUALITY AND INFERENCE SPEED ACROSS POSITIONAL ENCODING METHODS AND LANGUAGE PAIRS. EN-DE REPORTS MEAN $\pm$ STD OVER 2-3 SEEDS; HI-EN AND BN-EN USE ONE SEED. BOLD = BEST PER COLUMN. $\uparrow/\downarrow$ INDICATE PREFERRED DIRECTION.

Method	En-De			Hi-En			Bn-En			Decode s $\downarrow$
	BLEU $\uparrow$	chrF $\uparrow$	TER $\downarrow$	BLEU $\uparrow$	chrF $\uparrow$	TER $\downarrow$	BLEU $\uparrow$	chrF $\uparrow$	TER $\downarrow$	
RoPE	20.83 $\pm$ 0.03	52.51 $\pm$ 0.19	68.20 $\pm$ 0.06	7.78	38.29	158.74	17.87	42.70	79.56	62.0 $\pm$ 0.3
AdaptiveRoPE	20.69 $\pm$ 0.05	52.40 $\pm$ 0.02	68.45 $\pm$ 0.12	9.08	43.22	149.05	18.01	42.90	79.82	68.0 $\pm$ 0.4
GatesOnly	20.68 $\pm$ 0.08	52.45 $\pm$ 0.08	68.48 $\pm$ 0.19	8.37	41.43	169.07	17.79	42.97	79.57	67.6 $\pm$ 0.3
PhasesOnly	20.75 $\pm$ 0.05	52.46 $\pm$ 0.11	68.47 $\pm$ 0.03	9.42	43.26	175.90	17.06	42.16	83.16	71.6 $\pm$ 0.0
Sinusoidal	20.18 $\pm$ 0.15	52.01 $\pm$ 0.06	69.06 $\pm$ 0.08	—	—	—	17.40	40.58	75.05	54.7 $\pm$ 0.5
ALiBi	19.60	51.53	70.04	—	—	—	—	—	—	55.6




Fig. 2. Per-layer attention entropy (En-De). ALiBi maintains higher entropy across all layers, producing more uniform attention.

other methods (1.65–1.67), indicating more uniform attention distributions. Its higher sink ratio (0.114 vs. 0.078 for RoPE) reflects stronger focus on initial tokens — a known ALiBi property that stabilises long-sequence attention [4].

Sinusoidal also shows elevated entropy and sink ratio, consistent with its absolute encoding strategy. Among RoPE variants, entropy and sink values are nearly identical, confirming that the learned gates and phases do not fundamentally alter attention structure.

Figure 2 shows per-layer entropy. All methods exhibit a characteristic “rise-and-fall” pattern: entropy increases through the mid-layers (layers 1–2) as contextual representations form, then decreases in upper layers as attention sharpens. ALiBi maintains higher entropy at every layer, explaining its robustness to length variation.

TABLE III

BLEU BY SOURCE SENTENCE LENGTH BUCKET ON EN-DE (GREEDY DECODING, MEAN OVER SEEDS). TRAINING DISTRIBUTION SPANS BUCKETS 1–50 WORDS.

Method	1-10	11-20	21-30	31-50	51+
RoPE	18.51 $\pm$ 0.18	21.09 $\pm$ 0.16	20.49 $\pm$ 0.00	21.39 $\pm$ 0.00	22.59 $\pm$ 0.13
ARoPE	17.86 $\pm$ 0.90	20.86 $\pm$ 0.01	20.35 $\pm$ 0.12	21.48 $\pm$ 0.02	22.66 $\pm$ 0.94
Gates	18.45 $\pm$ 0.12	20.76 $\pm$ 0.27	20.28 $\pm$ 0.11	21.53 $\pm$ 0.14	22.19 $\pm$ 0.62
Phases	18.46 $\pm$ 0.27	20.99 $\pm$ 0.11	20.56 $\pm$ 0.16	21.22 $\pm$ 0.18	21.65 $\pm$ 0.18
Sinus	17.94 $\pm$ 0.36	20.12 $\pm$ 0.31	19.88 $\pm$ 0.02	21.07 $\pm$ 0.35	21.76 $\pm$ 0.42
ALiBi	17.77	19.63	19.28	20.23	22.23

VI. ANALYSIS OF LEARNED PARAMETERS

A. Learned Frequency Gates

Table VI reports gate statistics for En-De. All gates are below one, confirming systematic suppression of the RoPE frequency spectrum. Query gates (mean 0.915) are slightly higher than key gates (0.903), indicating more aggressive frequency suppression in keys.

B. Frequency Spectrum Shift

Figure 3 shows the effective frequency spectrum. The learned model downscales effective frequencies across all dimensions, with stronger suppression at higher-frequency indices. This means the model prefers *longer-range* positional signals, consistent with the need to align source and target with potentially different word order.




Fig. 3. Effective frequency spectrum: standard RoPE (dashed) vs. AdaptiveRoPE (solid). Shaded region = suppressed frequencies.

TABLE IV

BLEU DEGRADATION ON OOD SOURCE LENGTHS (EN-DE, GREEDY). VALUES ARE % CHANGE RELATIVE TO THE 1-20 WORD BUCKET.

Method	21-40	41-60	61-80	81-120	201-400
RoPE	+2.3	+6.7	−14.5	−36.7	−97.5
ARoPE	+0.1	+6.4	−15.8	−33.6	−97.9
Gates	+1.4	+8.1	−14.2	−31.8	−98.2
Phases	−0.2	+5.4	−14.4	−33.2	−98.1
Sinus	+5.1	+10.4	−11.9	−30.9	−95.8
ALiBi	+1.3	+6.1	+0.3	−17.0	−94.8

TABLE V

ATTENTION ENTROPY AND SINK RATIO BY PE METHOD (EN–DE, MEAN ACROSS SEEDS). HIGHER ENTROPY = MORE UNIFORM ATTENTION; HIGHER SINK = STRONGER FOCUS ON FIRST TOKEN.

Method	RoPE	ARoPE	Gates	Phases	Sinus	ALiBi
Entropy↑	1.658±0.004	1.664±0.001	1.658±0.018	1.649±0.011	1.673	2.233
Sink↑	0.078±0.004	0.078±0.004	0.080±0.002	0.077±0.002	0.108	0.114

TABLE VI

LEARNED GATE AND PHASE STATISTICS (EN–DE, SEED 43). INIT: GATES = 1.0, PHASES = 0.0.

Parameter	Mean	Std	Min	Max
g_q (query gates)	0.9151	0.0249	0.8479	0.9744
g_k (key gates)	0.9030	0.0285	0.7784	0.9631
ϕ_q (query phases)	0.0045	0.0099	−0.0216	0.0403
ϕ_k (key phases)	−0.0045	0.0099	−0.0404	0.0217

C. Per-Head and Per-Layer Specialisation

Figure 4 visualises gate values per head and frequency. Different heads learn distinct suppression profiles, confirming that uniform frequency weighting is a constraint the model would relax. Head 2 shows the smallest deviation (least suppression), while head 0 shows the largest, suggesting specialisation into long-range vs. short-range positional attention.

Figure 5 shows per-layer encoder gates. Earlier layers (0–2) suppress mid-to-high frequencies more aggressively, while later layers (4–5) approach 1.0. This suggests the encoder progressively relies more on positional information as representations become contextualised.




Fig. 4. Learned query gates g_q per head (rows) and frequency (columns). Red = suppression (< 1); blue = amplification (> 1).




Fig. 5. Per-layer encoder gates, averaged over heads. Earlier layers suppress more aggressively.

VII. DISCUSSION

Why are positional encoding requirements language-dependent? Our results suggest that typological properties — morphological complexity and word-order flexibility — determine which PE components matter. Hindi’s rich fusional morphology and relatively free word order benefit from *phase offsets*, which provide flexible positional shifts without altering the frequency spectrum. Bengali’s more regular agglutinative structure benefits from *frequency scaling*, which adjusts the positional “resolution” of different heads. German’s rigid SOV/V2 syntax is well-served by RoPE’s defaults, leaving little room for adaptation to improve.

The accuracy–extrapolation trade-off. RoPE’s rotary structure excels at in-domain lengths but fails catastrophically beyond the training distribution (−97% at 201–400 words). ALiBi’s linear bias sacrifices 1.2 BLEU in-domain for dramatically better OOD performance (−47%). This trade-off is fundamental: RoPE’s periodic encoding provides fine-grained relative position discrimination, but the periodicity creates aliasing at unseen lengths; ALiBi’s monotonic bias generalises linearly but lacks the fine-grained structure needed for peak in-domain performance.

Implications for MT systems. For high-resource language pairs with controlled input lengths (e.g., document translation with length filtering), RoPE remains the best choice. For low-resource or length-variable settings (e.g., speech-to-text, live translation), ALiBi or hybrid approaches may be preferable. Our cross-lingual findings suggest that future work should explore *language-conditional* positional encodings that select gate/phase profiles based on typological features.

Limitations. Experiments are at a single scale (47M parameters). While En–De uses 3 seeds, Hi–En and Bn–En use 1 seed due to computational constraints. We evaluate greedy decoding; beam search may reveal additional patterns. Whether findings generalise to larger models and more language pairs remains open.

VIII. CONCLUSION

We presented AdaptiveRoPE, a minimal learnable extension of RoPE that reveals how encoder-decoder MT models use the positional frequency spectrum. Through experiments on three language pairs (En–De, Hi–En, Bn–En) against five baselines, we established four key findings:

- 1) Trained gates consistently converge below one (mean ≈ 0.90), confirming a systematic low-frequency preference in encoder-decoder MT.

- 2) The optimal adaptation is language-dependent: phase offsets are critical for Hindi–En (+21% BLEU over RoPE), while frequency gates matter more for Bengali–En.
- 3) RoPE achieves the best in-domain BLEU, but ALiBi generalises better to OOD lengths (−47% vs. −97% degradation).
- 4) Attention entropy analysis reveals ALiBi’s higher entropy explains its length robustness.

These results challenge the “one-size-fits-all” assumption underlying current positional encoding practice and motivate future work on language-adaptive and task-adaptive PE designs.

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention Is All You Need,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [2] P. Shaw, J. Uszkoreit, and A. Vaswani, “Self-Attention with Relative Position Representations,” in *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, vol. 2, 2018, pp. 464–468.
- [3] C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu, “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer,” *Journal of Machine Learning Research*, vol. 21, no. 140, pp. 1–67, 2020.
- [4] O. Press, N. A. Smith, and M. Lewis, “Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation,” in *International Conference on Learning Representations*, 2022.
- [5] J. Su, Y. Lu, S. Pan, A. Murtadha, B. Wen, and Y. Liu, “RoFormer: Enhanced Transformer with Rotary Position Embedding,” *arXiv preprint arXiv:2104.09864*, 2021.
- [6] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, A. Rodriguez, A. Joulin, E. Grave, and G. Lample, “LLaMA: Open and Efficient Foundation Language Models,” *arXiv preprint arXiv:2302.13971*, 2023.
- [7] A. Q. Jiang, A. Sablayrolles, A. Mensch, C. Bamford, D. S. Chaplot, D. de las Casas, F. Bressand, G. Lengyel, G. Lample, L. Saulnier, L. Renard Lavaud, M.-A. Lachaux, P. Stock, T. Le Scao, T. Lavril, T. Wang, T. Lacroix, and W. El Sayed, “Mistral 7B,” *arXiv preprint arXiv:2310.06825*, 2023.
- [8] Z. Men, Y. Wu, X. Chen, and Z. Li, “Round and Round We Go! What Makes Rotary Positional Encodings Useful?” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.
- [9] B. Peng, J. Quesnelle, H. Fan, and E. Shippole, “YaRN: Efficient Context Window Extension of Large Language Models,” *arXiv preprint arXiv:2309.00071*, 2023.
- [10] Y. Ding, L. L. Zhang, C. Zhang, Y. Xu, N. Shang, J. Xu, F. Yang, and M. Yang, “LongRoPE: Extending LLM Context Window Beyond 2 Million Tokens,” in *International Conference on Machine Learning*, 2024.
- [11] O. Bojar, C. Buck, C. Federmann, B. Haddow, P. Koehn, J. Leveling, C. Monz, P. Pecina, M. Post, H. Saint-Amand, R. Soricut, L. Specia, and A. Tamchyna, “Findings of the 2014 Workshop on Statistical Machine Translation,” in *Proceedings of the Ninth Workshop on Statistical Machine Translation*, 2014, pp. 12–58.
- [12] G. Ramesh, S. Doddapaneni, A. Bheemaraj, M. Jobanputra, R. Akula, A. Sharma, S. Sahoo, H. Diddee, D. Kakwani, N. Kumar, S. Pradeep, S. Nagaraj, V. Deepak, K. Raghavan, A. Kunchukuttan, and P. Kumar, “Samanantar: The Largest Publicly Available Parallel Corpora Collection for 11 Indic Languages,” in *Transactions of the Association for Computational Linguistics*, vol. 10, 2022, pp. 145–162.
- [13] R. Dabre, H. Shrotriya, A. Kumar, and R. P. Bhatt, “IndicBART: A Pre-trained Model for Indian Languages,” in *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, 2021, pp. 3606–3614.
- [14] I. Loshchilov and F. Hutter, “Decoupled Weight Decay Regularization,” in *International Conference on Learning Representations*, 2019.
- [15] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, “BLEU: A Method for Automatic Evaluation of Machine Translation,” in *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, 2002, pp. 311–318.
- [16] M. Popović, “chrF: Character n-gram F-score for Automatic MT Evaluation,” in *Proceedings of the Tenth Workshop on Statistical Machine Translation*, 2015, pp. 392–395.
- [17] M. Post, “A Call for Clarity in Reporting BLEU Scores,” in *Proceedings of the Third Conference on Machine Translation: Research Papers*, 2018, pp. 186–191.